

Exercise 1:

Shortly answer the following questions:

- (a) What is the difference between inner and outer loss?
- (b) Which model is more likely to overfit the training data:
 - k-NN with 1 or with 10 neighbours?
 - Logistic regression with 10 or 20 features?
 - LDA or QDA?
- (c) Which of the following methods yield an unbiased generalization error estimate?
Performance estimation ...
 - on training data
 - on test data
 - on training and test data combined
 - using cross validation
 - using subsampling
- (d) Which problem does resampling of training and test data solve?
- (e) Which problem does nested resampling solve?

Exercise 2:

The Satellite dataset consists of pixels in 3x3 neighbourhoods in a satellite image, where each pixel is described by 4 spectral values, and the classification label of the central pixel. (for further information see [?Satellite](#)) We fit a k-NN model to predict the class of the middle pixel. The performance is evaluated with the mmce. Look at the following R code and output: The performance is estimated in different ways: using training data, test data and then with cross validation. How do the estimates differ and why? Which one should be used?

```
library(mlr3)
library(mlr3learners)
library(mlbench)

data(Satellite)
satellite_task <-
  TaskClassif$new(id = "satellite_task",
                  backend = Satellite,
                  target = "classes")
knn_learner <- lrn("classif.kknn", k = 3)

# Train and test subsets:
set.seed(42)
train_indices <-
  sample.int(nrow(Satellite), size = 0.8 * nrow(Satellite))
test_indices <- setdiff(1:nrow(Satellite), train_indices)
```

```

# Training data performance estimate
knn_learner$train(task = satellite_task, row_ids = train_indices)
pred <-
  knn_learner$predict(task = satellite_task, row_ids = train_indices)
pred$score()

## classif.ce
##          0

# Test data performance estimate
pred <-
  knn_learner$predict(task = satellite_task, row_ids = test_indices)
pred$score()

## classif.ce
## 0.09246309

# CV performance estimate
rdesc <- rsmp("cv", folds = 10)
res <- resample(satellite_task, knn_learner, rdesc)

## INFO [08:24:16.612] Applying learner 'classif.kknn' on task 'satellite_task' (iter 1/10)
## INFO [08:24:16.770] Applying learner 'classif.kknn' on task 'satellite_task' (iter 2/10)
## INFO [08:24:16.920] Applying learner 'classif.kknn' on task 'satellite_task' (iter 3/10)
## INFO [08:24:17.225] Applying learner 'classif.kknn' on task 'satellite_task' (iter 4/10)
## INFO [08:24:17.379] Applying learner 'classif.kknn' on task 'satellite_task' (iter 5/10)
## INFO [08:24:17.514] Applying learner 'classif.kknn' on task 'satellite_task' (iter 6/10)
## INFO [08:24:17.806] Applying learner 'classif.kknn' on task 'satellite_task' (iter 7/10)
## INFO [08:24:17.941] Applying learner 'classif.kknn' on task 'satellite_task' (iter 8/10)
## INFO [08:24:18.085] Applying learner 'classif.kknn' on task 'satellite_task' (iter 9/10)
## INFO [08:24:18.230] Applying learner 'classif.kknn' on task 'satellite_task' (iter 10/10)

res$score()

##           task           task_id           learner  learner_id
## 1: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 2: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 3: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 4: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 5: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 6: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 7: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 8: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 9: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
## 10: <TaskClassif> satellite_task <LearnerClassifKKNN> classif.kknn
##           resampling resampling_id iteration prediction classif.ce
## 1: <ResamplingCV>           cv           1      <list> 0.09627329
## 2: <ResamplingCV>           cv           2      <list> 0.07919255
## 3: <ResamplingCV>           cv           3      <list> 0.08540373
## 4: <ResamplingCV>           cv           4      <list> 0.09472050
## 5: <ResamplingCV>           cv           5      <list> 0.10559006
## 6: <ResamplingCV>           cv           6      <list> 0.08553655
## 7: <ResamplingCV>           cv           7      <list> 0.08864697
## 8: <ResamplingCV>           cv           8      <list> 0.10108865
## 9: <ResamplingCV>           cv           9      <list> 0.08864697
## 10: <ResamplingCV>          cv          10      <list> 0.10419907

```

```
res$aggregate()

## classif.ce
## 0.09292983
```

Exercise 3:

In preparing this course you already learned about `mlr3`. If you need to refresh your knowledge you can find help at <https://mlr3book.mlr-org.com/> under 'Basics'.

- How many performance measures do you already know? Try to explain some of them. How can you see which of them are available in `mlr3`?
- Use the `boston.housing` regression task from `mlr3` and split the data into 50 % training data and 50 % test data while training and predicting (i. e., use the `row_ids` argument of the `train` and `predict` function). Fit a prediction model (e. g. k-NN) to the training set and make predictions for the test set.
- Compare the performance on training and test data. Use the `score` function.
- Now use different observations (but still 50 % of them) for the training set. How does this affect the predictions and the error estimates of the test data?
- Use 10 fold cross-validation to estimate the performance. Hint: Use the `mlr` functions `rsmp` and `resample`.

Exercise 4:

Given are the results of a scoring algorithm and the associated *true* classes of 10 observations:

ID	Actual Class	Score
1	0	0.33
2	0	0.27
3	1	0.11
4	1	0.38
5	1	0.17
6	0	0.63
7	1	0.62
8	1	0.33
9	0	0.15
10	0	0.57

- Create a confusion matrix assuming the decision boundary at 0.5.
- Calculate: precision, sensitivity, negative predictive value, specificity, accuracy, error rate and F-measure.
- Draw the ROC curve and interpret it. Feel free to use R for the drawing.
- Calculate the AUC.