

Introduction to Machine Learning

Clustering

Bernd Bischl, Christoph Molnar, Daniel Schalk, Fabian Scheipl

Department of Statistics – LMU Munich



Hierarchical Clustering

MOTIVATION FOR CLUSTERING

Consider multivariate data with N observations (e.g. customers) and P features (e.g. characteristics of customers).

Task: divide data into groups (clusters), such that

- the observations in each cluster are as "similar" as possible (homogeneity within each cluster), and
- the clusters are as "far away" as possible from other clusters (heterogeneity between different clusters).

CLUSTERING VS. CLASSIFICATION

- In classification, the groups are known and we try to learn what differentiates these groups (i.e., learn a classification function) to properly classify future data.
- In clustering, we look at data, where groups are unknown and try to find similar groups.

Why do we need clustering?

- Discovery: looking for new insights in the data (e.g. finding groups of customers that buy a similar product).
- Derive a reduced representation of the full data set.

HIERARCHICAL CLUSTERING

Hierarchical clustering is a recursive process that builds a hierarchy of clusters. We distinguish between:

- ❶ Agglomerative (or bottom-up) clustering:
 - Start: Each observations is an *individual cluster*.
 - Repeat: Merge the two closest clusters.
 - Stop when there is only one cluster left.
- ❷ Divisive (or top-down) clustering:
 - Start: All observations are within *one* cluster.
 - Repeat: Divide the cluster that results in two clusters with biggest distance.
 - Stop when each observation is an individual cluster.

HIERARCHICAL CLUSTERING

—FIX THE TABLE—

Let X_1, \dots, X_N be observations with P features (dimensions), where $X_i = (x_{i1}, \dots, x_{iP})^\top$. A data set is a $(N \times P)$ -matrix of the form:

| | feature 1 | ... | ... | feature P |

| : — : : — : : — : : — : : — : | $|X_1|$ | x_{11} | ... | ... | x_{1P} | | :

: : : : | X_N | x_{N1} | ... | ... | x_{NP} |

HIERARCHICAL CLUSTERING

For hierarchical clustering, we need a definition for

- distances $d(X_i, X_j)$ between two observations X_i and X_j :
 - manhattan distance:


$$d(X_i, X_j) = ||X_i - X_j||_1 = \sum_{k=1}^P |x_{ik} - x_{jk}|$$

- euclidean distance:

$$d(X_i, X_j) = ||X_i - X_j||_2 = \sqrt{\sum_{k=1}^P (x_{ik} - x_{jk})^2}$$

- distances between two clusters (called linkage).

DISTANCES BETWEEN OBSERVATIONS



- manhattan: sum up the absolute distances in each dimension. In R: `"dist(data, method = "manhattan")"`
- euclidean: remember Pythagoras theorem from school? In R: `"dist(data, method = "euclidean")"`
- gower: can be used for mixed variables (categorical and numeric). In R: `"gower_dist()"` from the `"gower"` package
- see `"?dist"` for other distances.

GOWER DISTANCE I

- The Gower's metric calculates the distance between observations X_i and X_j for each feature separately and based on its data type (i.e., categorical or numeric).
- For a categorical feature X_k , the distance between the i -th and the j -th observation X_{ik} and X_{jk} is defined by

$$s_{ijk} = \begin{cases} 0 & \text{if } X_{ik} = X_{jk} \\ 1 & \text{if } X_{ik} \neq X_{jk}. \end{cases}$$

- For a numerical feature X_k , the distance between the i -th and the j -th observation X_{ik} and X_{jk} is defined by

$$s_{ijk} = \frac{|X_{ik} - X_{jk}|}{\max(X_k) - \min(X_k)}, \quad \text{so that } 0 \leq s_{ijk} \leq 1.$$

GOWER DISTANCE II

The Gower's metric S combines all individual distances of each feature by

$$S_{ij} = \frac{\sum_{k=1}^P w_k s_{ijk}}{\sum_{k=1}^P w_k},$$

where

- P : number of features.
- w_k : weight for feature k (typically $w_k = 1$).
- s_{ijk} : the difference (distance) between X_{ik} and X_{jk} , i.e. the i -th and j -th observation of feature k .

GOWER DISTANCE III - EXAMPLE

Observation	X_1	X_2
1	4	a
2	3	b
3	6	a

$$w_k = 1 \text{ for } k \in \{1, 2\}$$

$$s_{121} = \frac{|4 - 3|}{6 - 3} = \frac{1}{3}$$

$$s_{122} = 1, \text{ since } a \neq b$$

$$S_{12} = \frac{\frac{1}{3} + 1}{2} = \frac{2}{3}$$

$$s_{131} = \frac{|4 - 6|}{6 - 3} = \frac{2}{3}$$

$$s_{132} = 0, \text{ since } a = a$$

$$S_{13} = \frac{\frac{2}{3} + 0}{2} = \frac{1}{3}$$



$$s_{231} = \frac{|3 - 6|}{6 - 3} = 1$$

$$s_{232} = 1, \text{ since } b \neq a$$

$$S_{23} = \frac{1 + 1}{2} = 1$$

DISTANCES BETWEEN OBSERVATIONS

It is often a good idea to *normalize* the data before computing distances, especially when the scale of features is different, e.g.:




On the right plot, the distance is dominated by 'shoe size'.

DISTANCES BETWEEN OBSERVATIONS

The normalized feature $\tilde{X}_{\text{height}}$ is computed using X_{height} by \leftarrow
Normalization of the height feature means: \rightarrow

$$\tilde{X}_{\text{height}} = \frac{X_{\text{height}} - \text{mean}(X_{\text{height}})}{\text{sd}(X_{\text{height}})}.$$

Distances based on normalized data are better comparable and robust in terms of linear transformations (e.g. unit conversion).




DISTANCES BETWEEN CLUSTERS (LINKAGE)

- Assume that all observations X_1, \dots, X_N belong to $K < N$ different clusters.
- The linkage of two clusters C_r and C_s is a "score" describing their distance.

The most popular and simplest linkages are

- *Single Linkage*
- *Complete Linkage**
- *Average Linkage**
- *Centroid Linkage*


SINGLE LINKAGE



Single linkage defines the distance of the *closest point pairs* from different clusters as the distance between two clusters:

$$d_{\text{single}}(C_r, C_s) = \min_{i \in C_r, j \in C_s} d(X_i, X_j)$$


COMPLETE LINKAGE



Complete linkage defines the distance of the *furthest point pairs* of different clusters as the distance between two clusters:

$$d_{\text{complete}}(C_r, C_s) = \max_{i \in C_r, j \in C_s} d(X_i, X_j)$$


AVERAGE LINKAGE



(Note: Plot only shows distances between all green points and *one* red point)

In average linkage, the distance between two clusters is defined as the average distance across *all* pairs of two different clusters.

CENTROID LINKAGE



Centroid linkage defines the distance between two clusters as the distance between the two cluster centroids. The centroid of a cluster C_s with N_s points is the mean value of each dimension:


$$\bar{X}_s = \frac{1}{N_s} \sum_{i \in C_s} x_i$$

EXAMPLE: HIERARCHICAL CLUSTERING

Agglomerative hierarchical clustering starts with all points forming their own cluster and iteratively merges them until all points form a single cluster containing all points.

Example:

Step 1: $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$




EXAMPLE: HIERARCHICAL CLUSTERING

Agglomerative hierarchical clustering starts with all points forming their own cluster and iteratively merges them until all points form a single cluster containing all points.

Example:

Step 1: $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$

Step 2: $\{1\}, \{2, 3\}, \{4\}, \{5\}$



EXAMPLE: HIERARCHICAL CLUSTERING


Agglomerative hierarchical clustering starts with all points forming their own cluster and iteratively merges them until all points form a single cluster containing all points.

Example:

Step 1: $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$

Step 2: $\{1\}, \{2, 3\}, \{4\}, \{5\}$

Step 3: $\{1\}, \{2, 3\}, \{4, 5\}$



EXAMPLE: HIERARCHICAL CLUSTERING

Agglomerative hierarchical clustering starts with all points forming their own cluster and iteratively merges them until all points form a single cluster containing all points.


Example:

Step 1: $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$

Step 2: $\{1\}, \{2, 3\}, \{4\}, \{5\}$

Step 3: $\{1\}, \{2, 3\}, \{4, 5\}$

Step 4: $\{1, 2, 3\}, \{4, 5\}$



EXAMPLE: HIERARCHICAL CLUSTERING

Agglomerative hierarchical clustering starts with all points forming their own cluster and iteratively merges them until all points form a single cluster containing all points.

Example:


Step 1: $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$

Step 2: $\{1\}, \{2, 3\}, \{4\}, \{5\}$

Step 3: $\{1\}, \{2, 3\}, \{4, 5\}$



Step 4: $\{1, 2, 3\}, \{4, 5\}$

Step 5: $\{1, 2, 3, 4, 5\}$



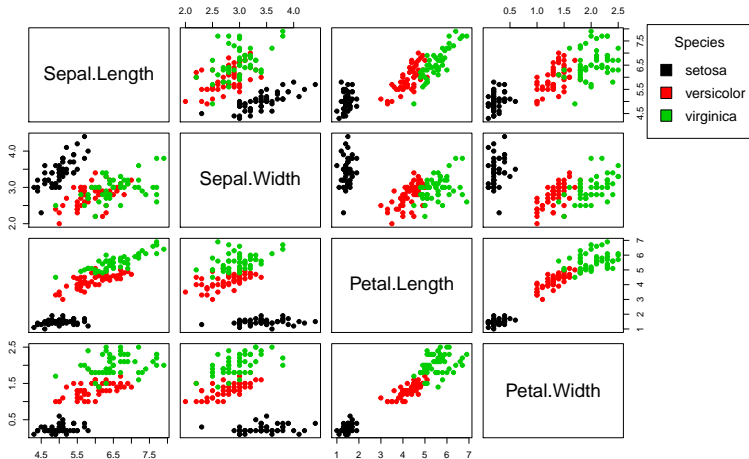
DENDROGRAM

A Dendrogram is a tree showing which clusters / observations are merged after each step. The 'height' is proportional to the distance between the two merged clusters:




R EXAMPLE WITH IRIS DATA

The data contains 150 leaf measurements for 3 flower species:




R EXAMPLE WITH IRIS DATA

We now "forget" the real groups specified by the 'Species' variable and try to find clusters based on the leaf measurements.

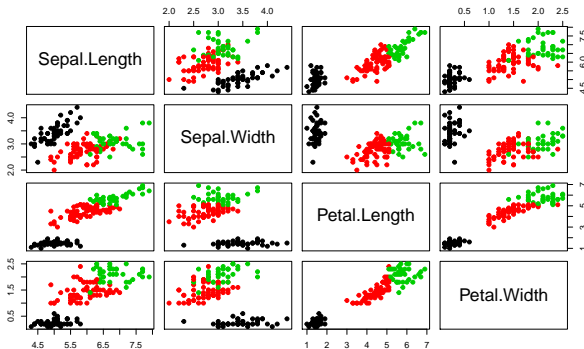


R EXAMPLE WITH IRIS DATA



R EXAMPLE WITH IRIS DATA



We can extract the clustering assignments by cutting the dendrogram, e.g. using $k = 3$ clusters:



PROPERTIES: SINGLE LINKAGE

Single linkage introduces the *chaining problem*:



- Only one pair of points needs to be close to merge clusters.
- A chain of points can expand a cluster over long distances.
- Points within a cluster can be too widely spread and not dense enough.



PROPERTIES: COMPLETE LINKAGE

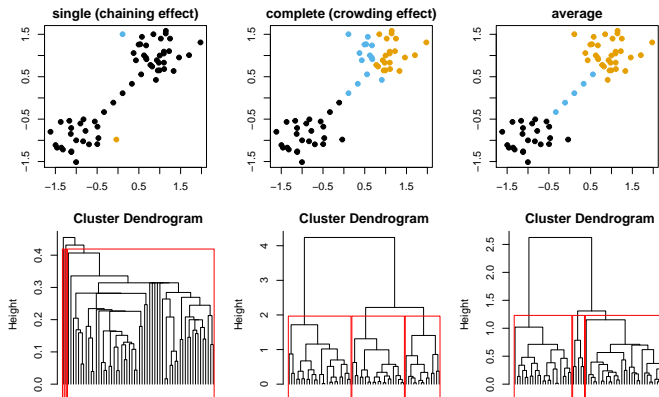
Complete linkage avoids chaining, but suffers from *crowding*:

- Merging is based on the furthest distance of point pairs from different clusters.
- Points of two different clusters can thus be closer than points within a cluster.
- Clusters are dense, but too close to each other. <!-- and sensitive to outliers. -->




PROPERTIES: AVERAGE LINKAGE

- Average linkage is based on the average distance between clusters and tries to avoid crowding and chaining.
- Produces clusters that are quite dense and rather far apart.



PROPERTIES: CENTROID LINKAGE

- Centroid linkage defines the distance based on **artificial data points** (the cluster centers), which produces dendrograms **with inversions**, i.e., the distance between the clusters to be merged can be smaller in the next step.
- In single, complete and average linkage, the distance between the clusters to be merged increases in each step. \Rightarrow always produces dendrograms **without inversions**.



SUMMARY

- *Hierarchical agglomerative clustering methods* iteratively merge observations/clusters until all observations are in one single cluster.
- Results in a hierarchy of clustering assignments which can be visualized in a *dendrogram*. Each node of the dendrogram represents a cluster and its 'height' is proportional to the distance of its child nodes.
- The most common linkage functions are *single*, *complete*, *average* and *centroid* linkage. There is no perfect linkage and each linkage has its own advantages and disadvantages.