LMU LUDWIG-MAXIMILIANS-UNIVERSITÄT MÜNCHEN

Introduction to Machine Learning

Chapter 2: Machine Learning Tasks

Bernd Bischl, Christoph Molnar, Daniel Schalk, Fabian Scheipl

Department of Statistics – LMU Munich

SUPERVISED LEARNING

- One tries to learn the relationship between “input” x and “output” y.
- For learning, there is training data with labels available
- Mathematically, we face a problem of function approximation: search for an f, such that, for all points in the training data, and also all newly observed points,

$$y \approx f(x).$$

SUPERVISED LEARNING

Regression Task

Goal: Predict a continuous output

- y is metric variable (with values in $\mathbb{R}$)
- Regression model can be constructed by different methods (e.g. trees or splines), not only statistical (linear) regression!

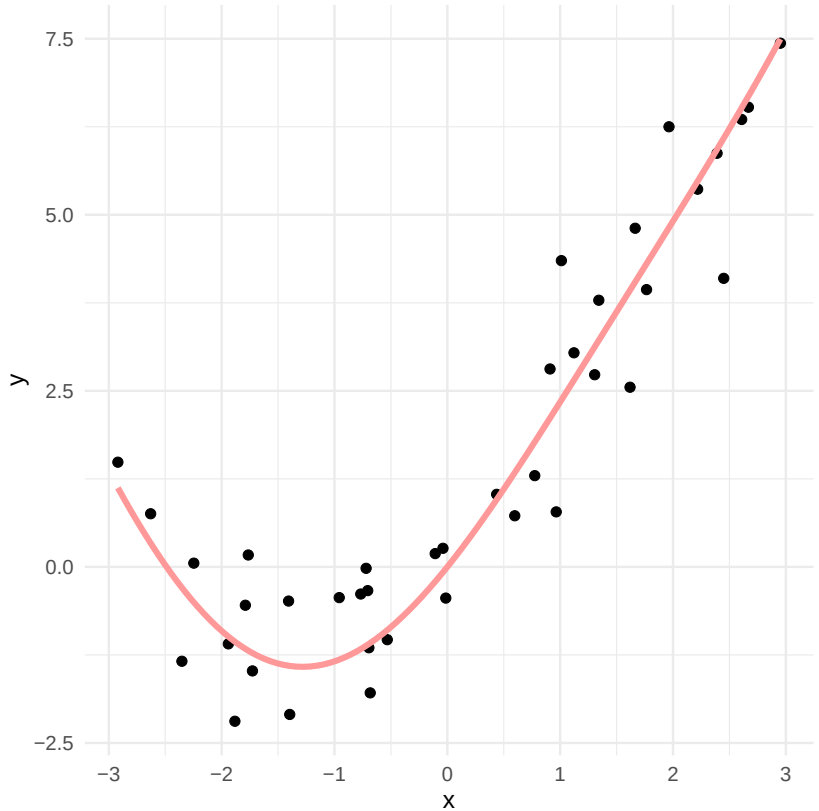

SUPERVISED LEARNING

Regression Task - Examples

- **Stock Trading:** Predicting the exact stock prices on the basis of company data and insider information
- **Pricing:** Anticipating the willingness-to-pay of new customers on the basis of purchases of other customers
- **Medicine:** Calculating the life expectancy for patients with a particular disease and severity (although life time analysis is often better here due to right censoring)
- **Income:** Predicting future income of a person based on education and skills

SUPERVISED LEARNING

Regression Task - Income Prediction

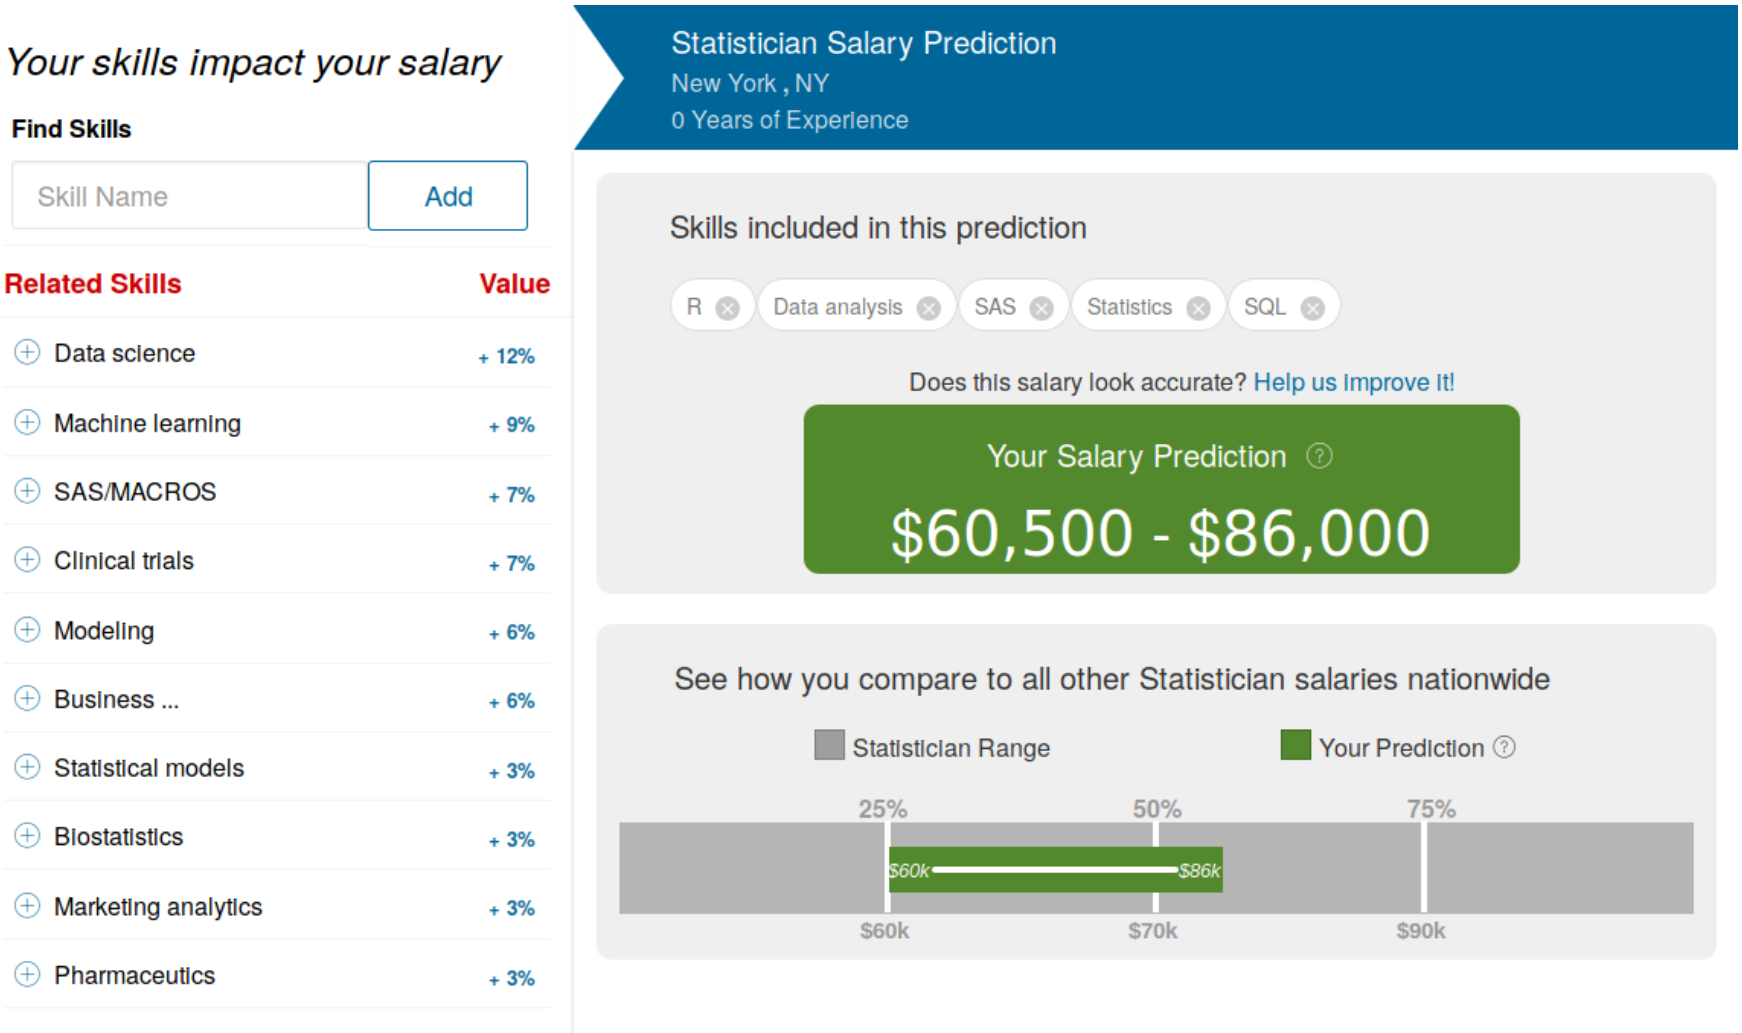

https://www.dice.com/salary-calculator

SUPERVISED LEARNING

Binary Classification Task

Goal: Predict a class (or membership probabilities)

- *y* is a categorical variable with two possible values
- Each observation belongs to exactly one class

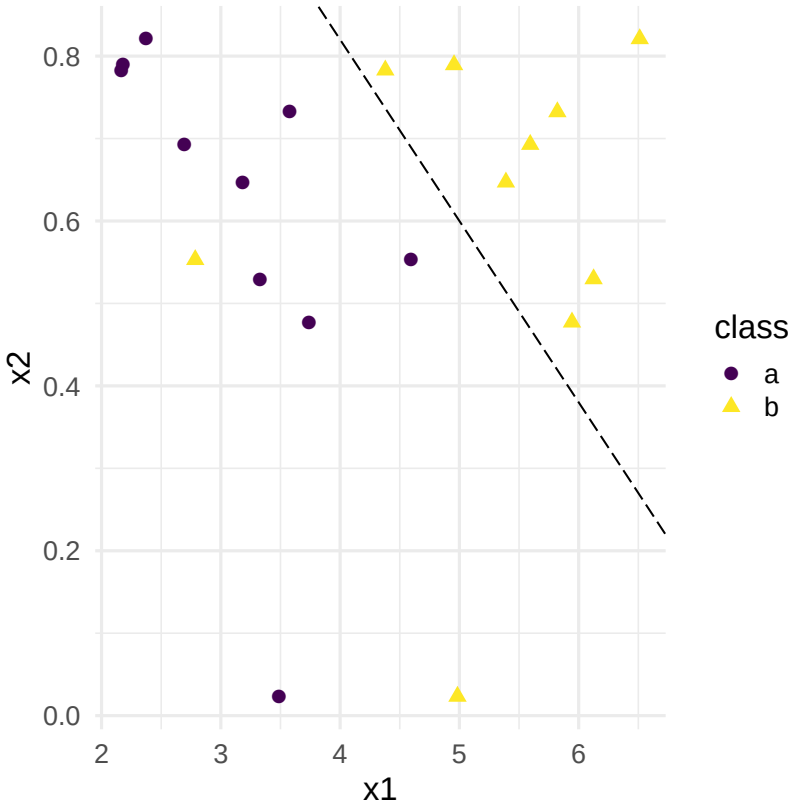

SUPERVISED LEARNING

Binary Classification Task - Examples

- **Credits:** Predicting credit fraud or default risk based on transactions
- **Medical Diagnosis:** Medically testing whether a patient has a specific illness or not
- **Software:** Detecting whether an e-mail is spam or not by using its content
- **Lie Detection:** Determine truthfulness of statements from physiological cues

SUPERVISED LEARNING

Binary Classification Task - Lie Detection

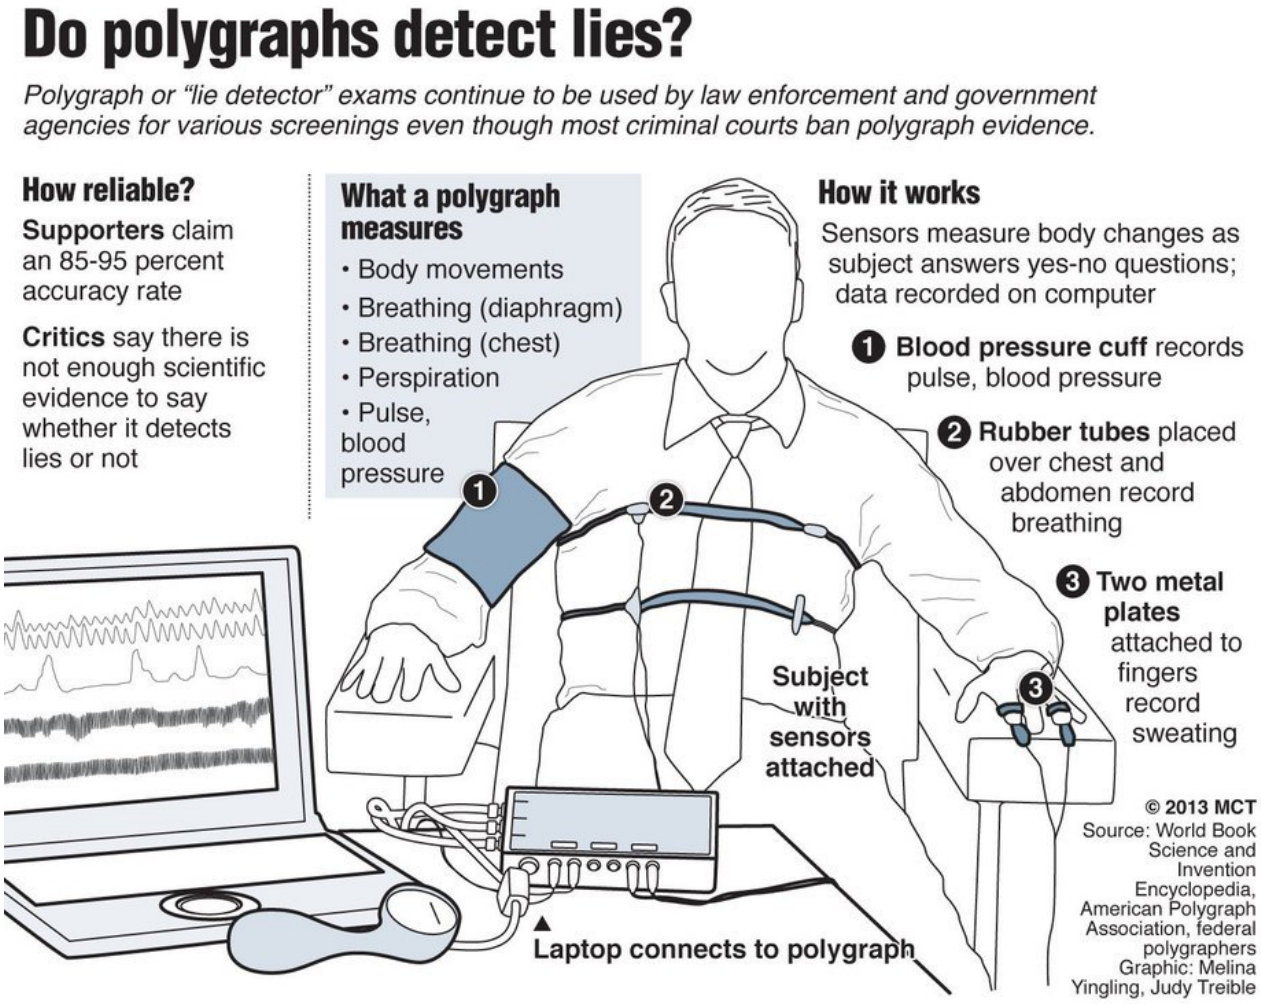

https://www.bendbulletin.com/localstate/deschutescounty/3430324-151/fact-or-fiction-polygraphs-just-an-investigative-tool

SUPERVISED LEARNING

Multiclass Classification Task

Goal: Predict a class (or membership probabilities)

- y is a categorical variable with more than two different unordered discrete values
- Each observation belongs to exactly one class

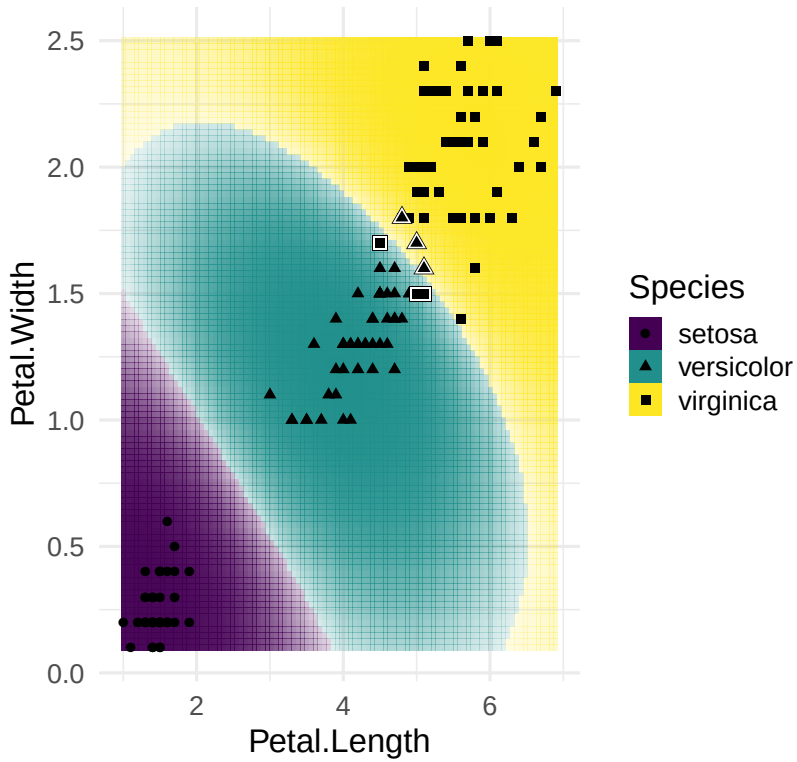

SUPERVISED LEARNING

Multiclass Classification Task - Examples

- **Image Recognition:** Deciding what animal (for example) a picture is showing
- **Stock Trading:** Identifying the best strategy for a specific stock (buy, sell, or wait) based on past prices
- **Biology:** Classifying plants and animals based on their exterior characteristics (e. g. iris flowers)
- **Medical Diagnosis:** Predicting a patients illness using the their symptoms

SUPERVISED LEARNING

Multiclass Classification Task - Medical Diagnosis

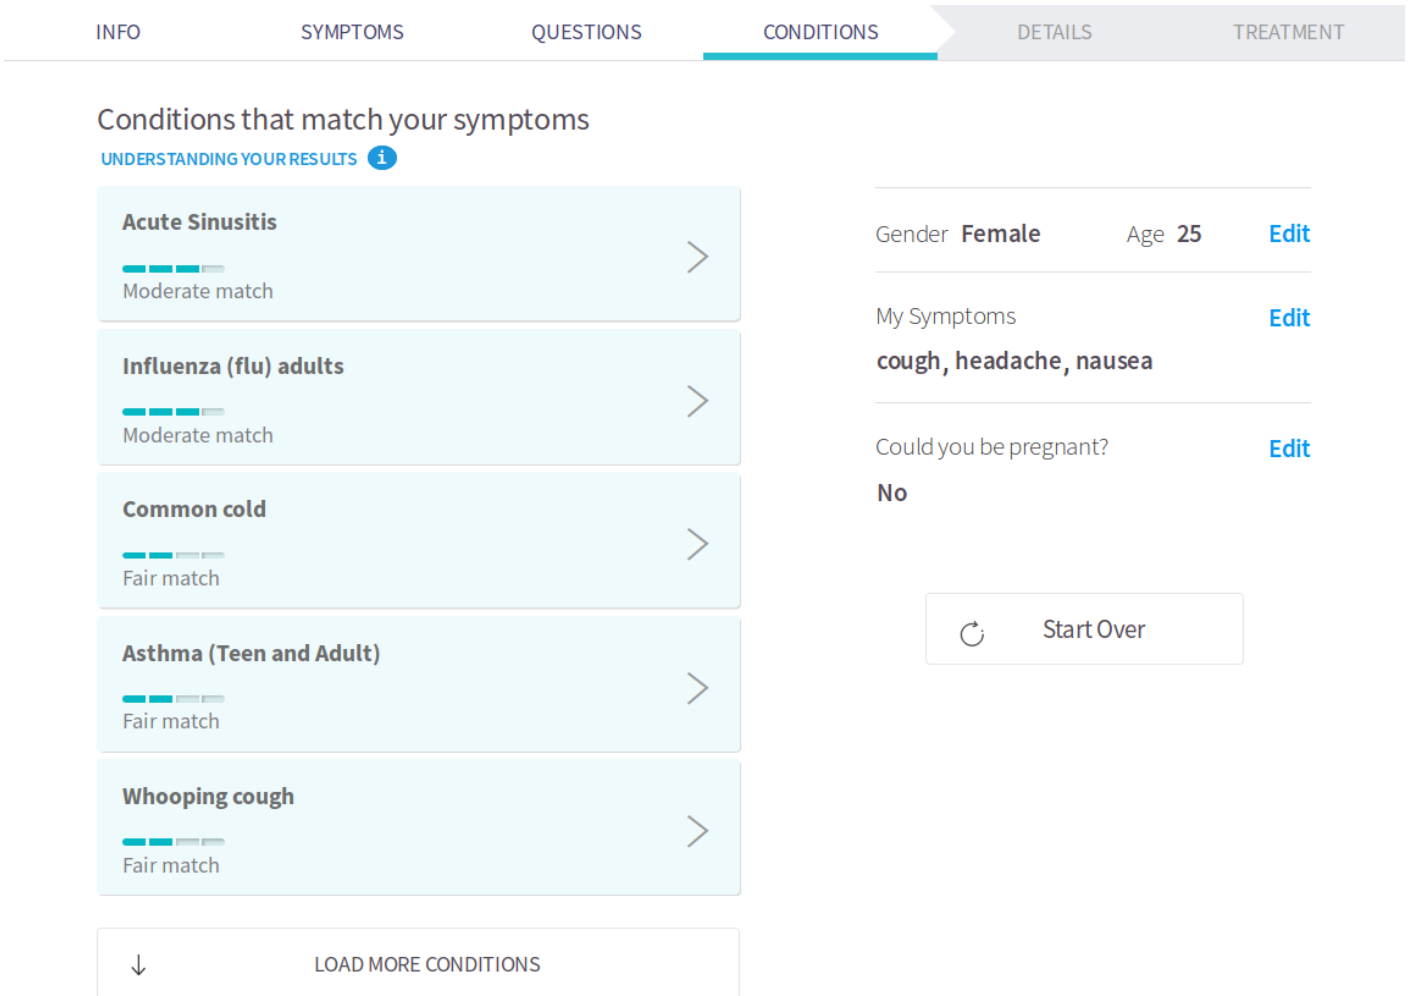

https://symptoms.webmd.com

SUPERVISED LEARNING

Classification Models

- Most classification models yield scoring functions for each of the g classes: $f(x) = (f_1(x), \ldots, f_g(x)) \in \mathbb{R}^g$.
- These are often called **discriminant functions**, their outputs are class scores or class probabilities.
- The actual classification rule is usually defined as:
  $h(x) = \underset{k \in \{1,\ldots,g\}}{\arg\max}\ f_k(x)$

SUPERVISED LEARNING

Other supervised learning tasks

- Multilabel classification
- Forecasting
- Survival prediction
- Cost-sensitive classification

ADDITIONAL LEARNING TASKS

Unsupervised learning

- Data without labels y
- Search for patterns within the inputs x
- *unsupervised* as there is no external criterion to optimize or “true” output
- Possible applications:
  - Dimensionality reduction (PCA, Autoencoders ...) : Compress information in $\mathcal{X}$
  - Clustering: Grouping similar observations, separating dissimilar observations
  - Outlier detection
  - Association rules

ADDITIONAL LEARNING TASKS

Semi-Supervised learning

- Large amount of labeled data necessary to train reliable model
- Creating labeled datasets often very expensive
- Learn from labeled (expensive) **and** unlabeled (cheap) data
- Unlabeled data in conjunction with a small amount of labeled data improves learning accuracy

Reinforcement learning

- Select actions in subsequent states within a certain environment to maximize lagged future reward
- Example: train neural net to play mario kart (environment)
  - Accelerate/ steer/ break (actions) at each time point (states) during playing
  - Reward: ranking after finish, should be maximized

MACHINE LEARNING TASKS

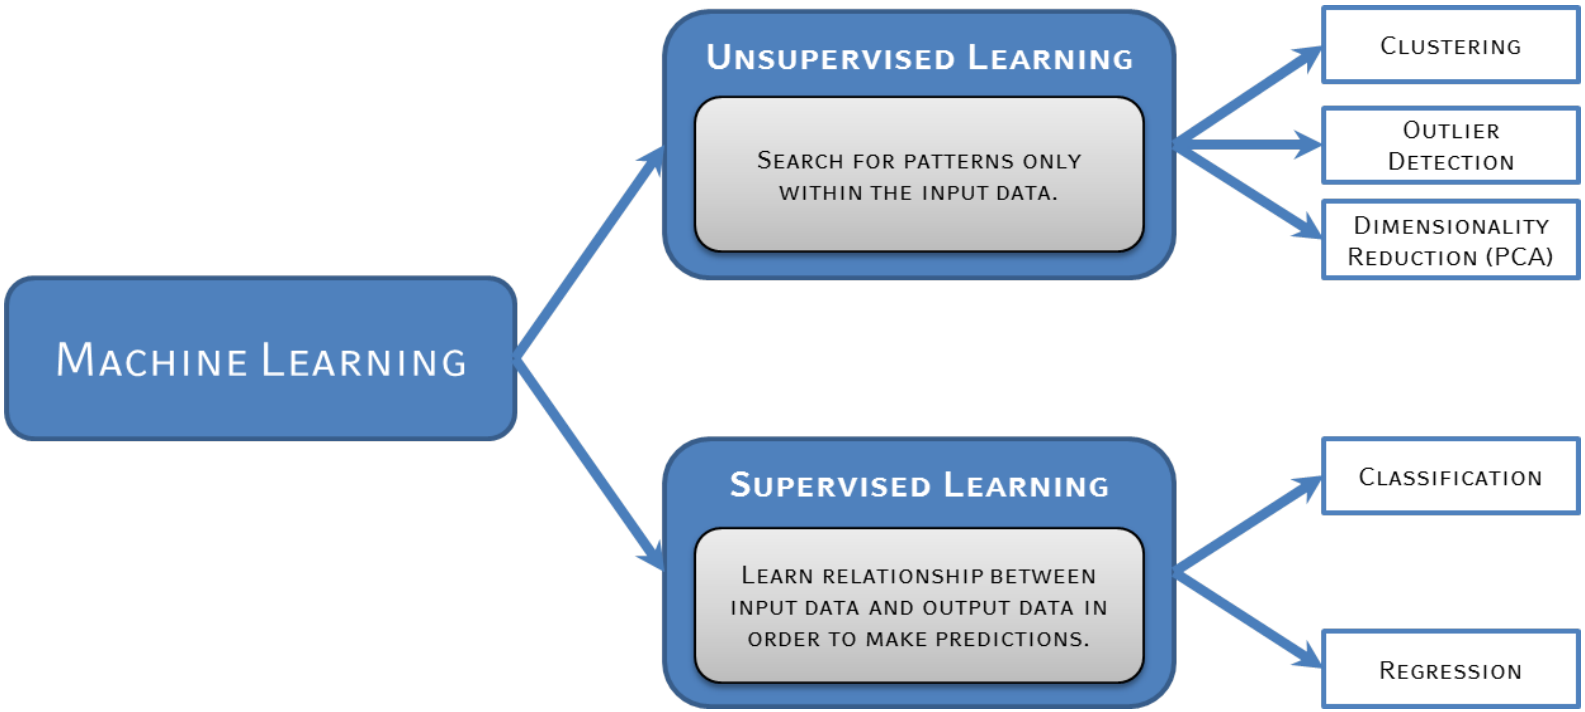

- In this course, we will deal with **supervised learning** for regression and classification only: predicting y based on x, using a model $f(x)$ that we learned from labeled training data.
- Classification models come with a slight twist: they typically learn g discriminant functions, and then these are turned into discrete predictions (details later).